



Python Data Analytics

OBA 410/510

Lundquist College of Business



UNIVERSITY OF OREGON
Lundquist College of Business

Two more Classification Techniques



Classification

- Regression: Target variable (y) is numerical
- Classification: Target variable (y) is binary or categorical
- Examples:
 - Is an email legitimate or spam?
 - Will a flight be on time or late?
 - Will a student receive an High Pass, Pass, Low Pass or Fail?
 - Is it an appropriate time to buy or sell a stock?
 - Will a customer choose to buy or not buy?
- The prediction in many Classification techniques is the probability of categories (or labels) for given set of input variables (X)



Default data

- Default of individual credit card payment ; 10,000 records
- Target : default (y)
 - Yes, No
- Features (Predictors) (X)
 - student (Yes(1), No(0)), balance, income

	default	student	balance	income
0	No	0	729.526495	44361.625070
1	No	1	817.180407	12106.134700
2	No	0	1073.549164	31767.138950
3	No	0	529.250605	35704.493940
4	No	0	785.655883	38463.495880
5	No	1	919.588530	7491.558572
6	No	0	825.513330	24905.226580
7	No	1	808.667504	17600.451340
8	No	0	1161.057854	37468.529290
9	No	0	0.000000	29275.268290



Descriptive Analysis

```
no=df_default['default']=='No'  
yes=df_default['default']=='Yes'
```

- Individuals who default tend to have **higher** balance
- Individuals who default tend to have **lower** income

```
df_default[no].describe()
```

	student	balance	income
count	9667.000000	9667.000000	9667.000000
mean	0.291404	803.943750	33566.166625
std	0.454433	456.476236	13318.251249
min	0.000000	0.000000	771.967729
25%	0.000000	465.714646	21405.060665
50%	0.000000	802.857102	34589.488060


```
df_default[yes].describe()
```

	student	balance	income
count	333.000000	333.000000	333.000000
mean	0.381381	1747.821690	32089.147124
std	0.486457	341.266808	13804.221110
min	0.000000	652.397134	9663.788159
25%	0.000000	1511.610952	19027.508630
50%	0.000000	1789.093391	31515.344490



Simple Logistic Regression

- Let's consider probability that an individual will default depends only on "balance"



$$\Pr(\text{default} = \text{Yes}) = \frac{e^{\beta_0 + \beta_1 \text{ balance}}}{1 + e^{\beta_0 + \beta_1 \text{ balance}}}$$

(Logistic function)

$$\Pr(\text{default} = \text{No}) = \frac{1}{1 + e^{\beta_0 + \beta_1 \text{ balance}}}$$

$$\Pr(\text{default} = \text{Yes}) + \Pr(\text{default} = \text{No}) = 1$$



Manipulation

$$\frac{\Pr(\text{default} = \text{Yes})}{1 - \Pr(\text{default} = \text{Yes})} = e^{\beta_0 + \beta_1 \text{ balance}}$$

- Left hand side of the above equation is call “odds”
- Odds close to **0** indicate **low** probability of default
- Odds close to **1** indicate **high** probability of default

$$\ln \left(\frac{\Pr(\text{default} = \text{Yes})}{1 - \Pr(\text{default} = \text{Yes})} \right) = \beta_0 + \beta_1 \text{ balance}$$

- Left hand side of the above example is call “log-odds”



Logistic Regression in scikit-learn

```
• X,y=df_default[['balance']],df_default['default']  
  from sklearn.model_selection import train_test_split  
  X_train, X_test, y_train, y_test=train_test_split(X,y,random_state=0)  
  
  from sklearn.linear_model import LogisticRegression  
  
  log_reg=LogisticRegression()  
  log_reg.fit(X_train,y_train)  
  
  print('Log reg acc on train: {:.3f}'.format(log_reg.score(X_train,y_train)))  
  print('Log reg acc on test: {:.3f}'.format(log_reg.score(X_test,y_test)))  
  
Log reg acc on train: 0.973  
Log reg acc on test: 0.968
```



```
p1=[1000]
p2=[2000]
p3=[3000]
print('Predcited traget for p1, p2, and p3:',log_reg.predict([p1,p2,p3]))
```

```
Predcited traget for p1, p2, and p3: ['No' 'No' 'Yes']
```

```
log_reg.predict_proba([p1,p2,p3])
```

```
array([[0.98723001, 0.01276999],
       [0.54594197, 0.45405803],
       [0.01835677, 0.98164323]])
```

```
log_reg.classes_
```

```
array(['No', 'Yes'], dtype=object)
```

```
print('Probablities for p1, p2, and p3:',log_reg.predict_proba([p1,p2,p3])[:,1])
```

```
Probablities for p1, p2, and p3: [0.01276999 0.45405803 0.98164323]
```



Logistic Regression

$$\text{Pr}(\text{target variable}) = \text{logistic}(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p)$$

$$\text{Pr}(\text{target variable}) = \frac{e^{(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p)}}{1 + e^{(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p)}}$$

$$\ln \left(\frac{\text{Pr}(\text{target variable})}{1 - \text{Pr}(\text{target variable})} \right) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_p x_p$$

- Coefficients (β_0, β_1, \dots) of the model are estimated using “Maximum Likelihood”

Logis

```
X,y=diabetes.iloc[:, :-1], diabetes['Outcome']  
X_train, X_test, y_train, y_test=train_test_split(X,y,random_state=0)
```

```
log_reg2=LogisticRegression()  
log_reg2.fit(X_train,y_train)
```

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
                    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
                    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
                    verbose=0, warm_start=False)
```

```
print('Log reg2 acc on train: {:.3f}'.format(log_reg2.score(X_train,y_train)))  
print('Log reg2 acc on test: {:.3f}'.format(log_reg2.score(X_test,y_test)))
```

```
Log reg2 acc on train: 0.757  
Log reg2 acc on test: 0.807
```

```
p1=[3,150,80,22,10,40,2.3,66]  
log_reg2.predict([p1])
```

```
array([1], dtype=int64)
```

```
log_reg2.predict_proba([p1])
```

```
array([[0.14143716, 0.85856284]])
```

```
log_reg2.classes_
```

```
array([0, 1], dtype=int64)
```

• Diabetes dataset



Logistic Regression

- By default, logistic regression applies an L2 regularization (in the same way that Ridge does for regression)
- L2 regularization, forces coefficients' magnitude to be closer to zero
- For ***LogisticRegression***, the trade-off parameter that determines the strength of the regularization is called **C**.
- Higher values of **C** correspond to less regularization
 - Higher C, LogisticRegression tries to fit the training set as best as possible
 - Lower C, more regularization, coefficients (β_i s) closer to zero

Log

```
cancer=pd.read_csv('breast_cancer_data.csv',index_col=0)
cancer.head()
```

UNIVERSITY OF OREGON
College of Business

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
id						
842302	M	17.99	10.38	122.80	1001.0	0.11840
842517	M	20.57	17.77	132.90	1326.0	0.16290

```
X,y=cancer.iloc[:,1:],cancer['diagnosis']
X_train,X_test,y_train,y_test=train_test_split(X,y,random_state=0)
```

```
log_reg3=LogisticRegression()
log_reg3.fit(X_train,y_train)
print('Log reg3 acc on train: {:.3f}'.format(log_reg3.score(X_train,y_train)))
print('Log reg3 acc on test: {:.3f}'.format(log_reg3.score(X_test,y_test)))
```

```
Log reg3 acc on train: 0.960
Log reg3 acc on test: 0.958
```



Logistic Regression in scikit-learn

```
: log_reg4=LogisticRegression(C=100)
log_reg4.fit(X_train,y_train)
print('Log reg4 acc on train: {:.3f}'.format(log_reg4.score(X_train,y_train)))
print('Log reg4 acc on test: {:.3f}'.format(log_reg4.score(X_test,y_test)))
```

```
Log reg4 acc on train: 0.967
Log reg4 acc on test: 0.965
```

```
: log_reg5=LogisticRegression(C=.1)
log_reg5.fit(X_train,y_train)
print('Log reg5 acc on train: {:.3f}'.format(log_reg5.score(X_train,y_train)))
print('Log reg5 acc on test: {:.3f}'.format(log_reg5.score(X_test,y_test)))
```

```
Log reg5 acc on train: 0.951
Log reg5 acc on test: 0.944
```




Transforming the data (scaling)

- Some techniques are sensitive to the scale of variables in data
 - Different scales in various variables
 - High variability in a single variable (**default** example with **balance** as the only feature)
- Therefore, before developing these models, it is a good idea to transform all the variables to the same scale
- One way to do this is using the *MinMaxScaler* function
- This function transforms all variables to a **between 0 and 1** scale
- It uses this formula:

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$



Transforming the data (scaling)

x1	x2
100	1.1
110	1
130	1.2
115	3
125	2.2

		x1	x1_new
max	130	100	0
min	100	110	0.3
		130	1
		115	0.5
		125	0.8

		x2	x2_new
max	3	1.1	0.05
min	1	1	0
		1.2	0.1
		3	1
		2.2	0.6

x1_new	x2_new
0.0	0.05
0.3	0
1.0	0.1
0.5	1
0.8	0.6



Evaluating the Effect of Parameters

- ***validation_curve*** function
 - Determines training and test scores for varying parameter values.
- Function inputs:
 - estimator: object type that implements the “fit” and “predict” methods
 - X: Training features
 - y: Target
 - param_name: Name of the parameter that will be varied.
 - param_range: The values of the parameter that will be evaluated.
 - cv : Determines the cross-validation splitting strategy

3 breast cancer data

```
: 1 # reading the data
   2 cancer=pd.read_csv('breast_cancer_data.csv',index_col=0)
   3 cancer.head(1)
```

```
:
      diagnosis radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean
id
842302      M      17.99      10.38      122.8      1001.0      0.1184

1 rows × 31 columns
```

```
: 1 X_cancer, y_cancer=cancer.iloc[:,1:], cancer['diagnosis']
   2 display(X_cancer.head(1))
   3 display(y_cancer.head(2))
```

```
      radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean
id
842302      17.99      10.38      122.8      1001.0      0.1184      0.277
```

1 rows × 30 columns

```
id
842302      M
842517      M
Name: diagnosis, dtype: object
```



```
3]: 1 from sklearn.preprocessing import MinMaxScaler
```

```
9]: 1 scaler1=MinMaxScaler()  
    2 X_cancer_trns=scaler1.fit_transform(X_cancer)  
    3 X_cancer_trns
```

```
9]: array([[0.52103744, 0.0226581 , 0.54598853, ..., 0.91202749, 0.59846245,  
          0.41886396],  
          [0.64314449, 0.27257355, 0.61578329, ..., 0.63917526, 0.23358959,  
          0.22287813],  
          [0.60149557, 0.3902604 , 0.59574321, ..., 0.83505155, 0.40370589,
```

```
2 from sklearn.model_selection import validation_curve
```

```
1 C_range=[.1,.5,1,10,100,200]
```

```
1 # validation_curve has two outputs: 1- score on training sets 2- scores on test sets  
2 # order is important  
3 train_scores, test_scores=validation_curve(LogisticRegression(solver='liblinear'),  
4                                             X_cancer, y_cancer,  
5                                             param_name='C',param_range=C_range, cv=4)
```

```
1 train_scores.round(4)
```

```
array([[0.9531, 0.9532, 0.9438, 0.9578],  
       [0.9577, 0.9555, 0.9508, 0.9672],  
       [0.9554, 0.9578, 0.9555, 0.9649],  
       [0.9695, 0.9696, 0.9649, 0.9696],  
       [0.9765, 0.9696, 0.9672, 0.9789],  
       [0.9742, 0.9742, 0.9789, 0.9742]])
```

```
1 print('ave cross val scores on train:',train_scores.mean(axis=1).round(4))
```

```
ave cross val scores on train: [0.952  0.9578 0.9584 0.9684 0.9731 0.9754]
```

```
1 test_scores.round(4)
```

```
array([[0.9021, 0.9648, 0.9577, 0.9085],  
       [0.9301, 0.9366, 0.9718, 0.9296],  
       [0.9301, 0.9366, 0.9718, 0.9366],  
       [0.9371, 0.9507, 0.9718, 0.9437],  
       [0.9301, 0.9507, 0.9718, 0.9507],  
       [0.9301, 0.9437, 0.9718, 0.9437]])
```

```
1 print('ave cross val scores on test:',test_scores.mean(axis=1).round(4))
```

```
ave cross val scores on test: [0.9333 0.942  0.9438 0.9508 0.9508 0.9473]
```




UNIVERSITY OF OREGON
Lundquist College of Business

Support Vector Machines



Support Vector Machine

- Approach for Classification (also for regression, but in this class we only use it for classification)
- Developed by Computer Scientists in 1990's
- In short stated as "SVM"
- There are two types of SVMs:
 - Linear SVM
 - Non-linear SVM




Terminology

- Hyperplane
 - In two dimension space: a line
 - In three dimension space: a plane
- Equation of the hyperplane in a p-dimensional space:
 - $\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p = 0$
- Example : A hyperplane in two-dimensional space $1 + 2X_1 + 3X_2 = 0$




Terminology





Classification using Separating Hyperplane


- Suppose we have training data of “n” observations and “p” variables
- Assume each of the observations fall into one of the two classes $\{-1, 1\}$
- Our goal is to classify a new test data, x^* with “p” variables into one of these classes
- All are hyperplanes shown in black are equally good





Maximum Margin Classifier

- Among all the hyperplanes, maximum margin classifier is the one that makes the biggest gap between the two classes
- The margin is the width that the decision boundary can be increased before hitting a data point.
- The objective in SVM
 - Minimizing the classification error and maximizing the margin between two classes





Support Vector Machine

Feature vector

Class value (target prediction)




- $f(X, \beta) = \text{sign}(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p)$

- $X=(-1, 0)$ what is the class?

- $X=(1, 1)$

$$1 + 2X_1 + 3X_2 = 0$$

Separating Hyperplane
(decision boundary)





Transforming the data (scaling) before using SVM

- SVM is very sensitive to the scale of variables in data
 - Different scales in various variables
 - High variability in a single variable (**default** example with **balance** as the only feature)
- Therefore, before developing SVM models, it is a good idea to transform all the variables to the same scale
- One way to do this is using the *MinMaxScaler* function
- This function transforms all variables to a **between 0 and 1** scale
- It uses this formula:

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$



Transforming the data (scaling) before using SVM

x1	x2
100	1.1
110	1
130	1.2
115	3
125	2.2

		x1	x1_new
max	130	100	0
min	100	110	0.3
		130	1
		115	0.5
		125	0.8

		x2	x2_new
max	3	1.1	0.05
min	1	1	0
		1.2	0.1
		3	1
		2.2	0.6

x1_new	x2_new
0.0	0.05
0.3	0
1.0	0.1
0.5	1
0.8	0.6



SVM in scikit-learn

```
: from sklearn.svm import LinearSVC
```

```
: df_default.head(n=2)
```

```
:
```

	default	student	balance	income
0	No	0	729.526495	44361.62507
1	No	1	817.180407	12106.13470

```
: X_def, y_def=df_default[['balance']],df_default['default']
```

```
91]: X_def.describe()
```

```
91]:
```

	balance
count	10000.000000
mean	835.374886
std	483.714985
min	0.000000
25%	481.731105
50%	823.636973
75%	1166.308387
max	2654.322576



SVM in scikit-learn

- LinearSVC, has its own random number generator
- If we do not specify the random_state option, every time we fit our model, we will have different scores

```
71]: from sklearn.preprocessing import MinMaxScaler
```

```
72]: scaler1=MinMaxScaler()
```

```
73]: X_def_trs=scaler1.fit_transform(X_def)  
X_def_trs
```

```
73]: array([[0.2748447 ],  
          [0.30786778],  
          [0.40445316],  
          ...,  
          [0.31850386],  
          [0.59111468],  
          [0.07569622]])
```


```
74]: X_train,X_test, y_train, y_test=train_test_split(X_def_trs,y_def,random_state=0)
```

```
92]: svm=LinearSVC(random_state=0)  
svm.fit(X_train,y_train)
```



SVM in scikit-learn

Test data points
have to be
transformed
before making
predictions



```
5]: svm.coef_  
6]: array([[3.01440675]])  
7]: svm.intercept_  
8]: array([-2.41896078])  
9]: # -2.418+3.014x
```

```
93]: print('svm on train: {:.2%}'.format(svm.score(X_train,y_train)))  
     print('svm on test: {:.2%}'.format(svm.score(X_test,y_test)))
```

```
svm on train: 97.11%  
svm on test: 96.60%
```

```
94]: p=[[800],[5000]]  
     p_trs=scaler1.transform(p)  
     p_trs
```

```
94]: array([[0.30139517],  
           [1.8837198 ]])
```

```
95]: svm.predict(p_trs)
```

```
95]: array(['No', 'Yes'], dtype=object)
```




Regularization Parameter

- Similar to LogisticRegression,
- By default, ***LinearSVC*** applies an L2 regularization (in the same way that Ridge does for regression)
- L2 regularization, forces coefficients' magnitude to be closer to zero
- For ***LinearSVC***, the trade-off parameter that determines the strength of the regularization is called ***C***.
- Higher values of ***C*** correspond to less regularization
 - Higher C, LogisticRegression tries to fit the training set as best as possible
 - Lower C, more regularization, coefficients (β_i s) closer to zero



SVM in scikit-learn

1.2 breast cancer data

```
|: 1 cancer=pd.read_csv('breast_cancer_data.csv', index_col=0)
   2 cancer.head(1)
```

```
|:
      diagnosis  radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  compactness_mean  conca
id
842302         M         17.99         10.38         122.8         1001.0         0.1184         0.2776
```

1 rows × 31 columns

```
|: 1 # create features and target sets
   2 X_cancer, y_cancer=cancer.iloc[:,1:], cancer['diagnosis']
   3 display(X_cancer.head(1))
   4 display(y_cancer.head(2))
```

```
      radius_mean  texture_mean  perimeter_mean  area_mean  smoothness_mean  compactness_mean  concavity_mean
id
842302         17.99         10.38         122.8         1001.0         0.1184         0.2776         0.3001
```

1 rows × 30 columns

```
id
842302    M
842517    M
Name: diagnosis, dtype: object
```



SVM in scikit-learn

```
|: 1 # transforming the features
2 from sklearn.preprocessing import MinMaxScaler
3 from sklearn.svm import LinearSVC
4 from sklearn.model_selection import validation_curve
```

```
|: 1 scaler2=MinMaxScaler()
2 X_cancer_trns=scaler2.fit_transform(X_cancer)
3 X_cancer_trns[:,:]
```

```
|: array([[0.52103744, 0.0226581 , 0.54598853, 0.36373277, 0.59375282,
          0.7920373 , 0.70313964, 0.73111332, 0.68636364, 0.60551811,
          0.35614702, 0.12046941, 0.3690336 , 0.27381126, 0.15929565,
          0.35139844, 0.13568182, 0.30062512, 0.31164518, 0.18304244,
          0.62077552, 0.14152452, 0.66831017, 0.45069799, 0.60113584,
          0.61929156, 0.56861022, 0.91202749, 0.59846245, 0.41886396]])
```

```
|: 1 C_range=[.1,1,5,10,50,100]
```

```
|: 1 train_scores, test_scores=validation_curve(LinearSVC(random_state=0, max_iter=100000)
2                                     ,X_cancer_trns,y_cancer,
3                                     param_name='C', param_range=C_range, cv=4)
```



SVM in scikit-learn

```
42]: 1 train_scores.round(4)
```

```
42]: array([[0.9624, 0.9578, 0.9672, 0.9696],  
          [0.9812, 0.9789, 0.9859, 0.9766],  
          [0.9883, 0.9859, 0.9883, 0.9813],  
          [0.9883, 0.9883, 0.9906, 0.9813],  
          [0.993 , 0.9906, 0.9906, 0.9836],  
          [0.9953, 0.9906, 0.9906, 0.9836]])
```

```
43]: 1 test_scores.round(4)
```

```
43]: array([[0.951 , 0.9437, 0.9577, 0.9718],  
          [0.972 , 0.9859, 0.9648, 0.9859],  
          [0.965 , 0.9789, 0.9648, 0.9859],  
          [0.972 , 0.9789, 0.9648, 0.9789],  
          [0.958 , 0.9789, 0.9859, 0.9648],  
          [0.951 , 0.9577, 0.9859, 0.9577]])
```

```
44]: 1 print('ave cross val scores on train:',train_scores.mean(axis=1).round(3))
```

```
ave cross val scores on train: [0.964 0.981 0.986 0.987 0.989 0.99 ]
```

```
45]: 1 print('ave cross val scores on test:',test_scores.mean(axis=1).round(3))
```

```
ave cross val scores on test: [0.956 0.977 0.974 0.974 0.972 0.963]
```



UNIVERSITY OF OREGON
Lundquist College of Business

Linear Models

parameters, strengths, weaknesses



Linear Models

parameters, strengths, weaknesses

- Linear regression, Lasso, Ridge
- Logistic regression, Linear support vector machine
- The main parameter is the regularization parameter
 - ***alpha*** in regression models and ***C*** in LinearSVC and LogisticRegression
 - Large values for ***alpha*** or small values for ***C*** mean simpler models
 - In particular for regression models, tuning this parameter is very important



Linear Models

parameters, strengths, weaknesses

- The other parameter in linear models is the regularization type
 - L1 versus L2
 - Default for LinearSVC, LogisticRegression and Ridge is L2
 - Lasso uses L1
 - If you have many features and think only a few of them are important, you should use L1, otherwise use the default (L2)
 - As L1 only uses a few features, it is easier to explain



Linear Models

parameters, strengths, weaknesses

- Linear models are very fast to train, also fast to predict
- They are suitable for very large and sparse datasets
- They perform well when the number of features is large compared to the number of samples (records)
- Linear models are usually easy to understand because of their linear nature!
- Linear models might not perform well in lower-dimensional spaces (fewer features in data)